



# KSFUPRO1K1U – Functional Programming

## Lecture 7: Imperative Features and Efficiency

Niels Hallenberg

These slides are based on original slides by Michael R. Hansen, DTU. Thanks!!!



The original slides has been used at a course in functional programming at DTU.



- Imperative programming,  
by simple examples.

# What is this?



```
let ...
  let rec visit u =
    color.[u] <- Gray ; time := !time + 1; d.[u] <- !time
    let rec h v = if color.[v] = White
                  then pi.[v] <- u
                   visit v
    List.iter h (adj.[u])
    color.[u] <- Black
    time      := !time + 1
    f.[u]     <- !time

  let mutable i = 0
  while i < V do
    if color.[i] = White
    then visit i
    i <- i + 1
  (d, f, pi);;
```



"Direct" translation of pseudocode from Corman, Leiserson, Rivest.

Remaining parts:

```
type color = White | Gray | Black;;

let dfs(V,adj: int list[]) =
  let color          = Array.create V White
  let pi             = Array.create V -1
  let d              = Array.create V -1
  let f              = Array.create V -1
  let time           = ref 0
  let rec visit u =
    ....

  let mutable i = 0
  while i < V do
    ....
    (d, f, pi);;
```




```
val (d,f,pi) = dfs(g6);
```

$d$  : Discovery times

$f$  : Finishing times

$pi$  : Predecessors

A node  $i$  is marked  $d_i/f_i$





A *store* is a table associating values  $v_i$  with locations  $l_i$ :

$$\left[ \begin{array}{ccc} l_1 & \mapsto & v_1 \\ l_2 & \mapsto & v_2 \\ & \dots & \\ l_n & \mapsto & v_n \end{array} \right]$$

## Allocation of a new cell in the store



```
let mutable x = 1;;  
val mutable x : int = 1
```

```
let mutable y = 3;;  
val mutable y : int = 3
```

Results in the following environment and store:

Environment

$$\left[ \begin{array}{l} x \mapsto l_1 \\ y \mapsto l_2 \end{array} \right]$$

Store

$$\left[ \begin{array}{l} l_1 \mapsto 1 \\ l_2 \mapsto 3 \end{array} \right]$$

A similar effect, but not exactly the same, is achieved by:

```
let x = ref 1;;  
let y = ref 3;;
```



Given the following environment and store:

Environment	Store
$\left[ \begin{array}{l} x \mapsto l_1 \\ y \mapsto l_2 \end{array} \right]$	$\left[ \begin{array}{l} l_1 \mapsto 1 \\ l_2 \mapsto 3 \end{array} \right]$

The assignment  $x \leftarrow y+2$  results in the new store:

$$\left[ \begin{array}{l} l_1 \mapsto 5 \\ l_2 \mapsto 3 \end{array} \right]$$

A similar effect, but not exactly the same, is achieved by the assignment  $x := !y + 2$

- The assignment  $x := \dots$  is used
- The explicit “contentsOf”  $!y$  is necessary

when  $\text{let } x = \text{ref } \dots$  and  $\text{let } y = \text{ref } \dots$  are used





You can also have a mutable record field

```
> type intRec = {mutable count : int } ;;
```

```
type intRec =  
  {mutable count: int;}
```

```
> let r1 = {count = 0} ;;
```

```
val r1 : intRec = {count = 0;}
```

```
> let incr x =  
  x.count <- x.count + 1  
  x.count ;;
```

```
val incr : intRec -> int
```

```
> incr r1;;  
val it : int = 1  
> incr r1;;  
val it : int = 2
```

Environment

$[ r1 \mapsto \{count \mapsto l_1\} ]$

Store

$[ l_1 \mapsto 2 ]$



Say we have a mutable record field and a few operators:

```
type 'a ref = {mutable contents: 'a}
let ref v = {contents = v}
let (!) r = r.contents
let (:=) r v = r.contents <- v
```

Then execute the following:

```
let a = ref 5
> val a : int ref = {contents = 5;}
let b = a
> val b : int ref = {contents = 5;}
```

A record is a reference type and hence heap allocated

Environment

$$\left[ \begin{array}{l} a \mapsto \{\text{contents} \rightarrow l_1\} \\ b \mapsto \{\text{contents} \rightarrow l_1\} \end{array} \right]$$

Store

$$[ l_1 \mapsto 5 ]$$

The two variables `a` and `b` are stack allocated and both pointing at the same record stored in the heap.

This is how references `ref`'s are implemented in F#.



One may ask why we have both `mutable` and `ref`. Other functional languages, e.g., Standard ML only have references.

- The compiler automatically inserts the `contentsOf` operator when you use `mutable`. You have explicitly to dereference a reference with `!`.
- A mutable value is stack allocated. A reference (`ref`) is heap allocated, as it is implemented as a record.
- You can alias two references. You can't do that with mutables (continuing from previous slide):

```

b := 10
> val it : unit = ()
!a
> val it : int = 10
!b
> val it : int = 10

let mutable a' = 5
> val mutable a' : int = 5

let mutable b' = a'
> val mutable b' : int = 5
b' <- 10
> val it : unit = ()
a'
> val it : int = 5
b'
> val it : int = 10

```

- Due to restrictions on polymorphic expressions, Section 4.5, you can't declare the following top level mutable: `let mutable a = [];;`.



Default values may be obtained using

```
Unchecked.defaultof<type>
```

```
> Unchecked.defaultof<int>;  
val it : int = 0
```

Two expressions are combined with  $; : exp_1; exp_2$

Evaluate first  $exp_1$  and then  $exp_2$ .

If  $exp_2$  has type  $\tau$ , then  $exp_1; exp_2$  has type  $\tau$  as well.

A warning is issued if  $exp_1$  has a type different from unit.

You may use `ignore` to overrule the warning

```
let mutable x = 0;;  
let y = (x;23);;
```

```
/Users/nielshallenberg/Dropbox/...: warning FS0020: This expression  
should have type 'unit', but has type 'int'. Use  
'ignore' to discard the result of the expression, or  
'let' to bind the result to a name.
```



- `"a []` is the type of one-dimensional, mutable, zero-based constant-time-access arrays with elements of type `'a`."

`Array.create n v` creates an array with `n` entries all containing `v`

Examples:


```
let a = Array.create 5 "a";;  
val a : string [] = [|"a"; "a"; "a"; "a"; "a"|]  
  
a.[2] <- "b";;  
val it : unit = ()  
  
a;;  
val it : string [] = [|"a"; "a"; "b"; "a"; "a"|]  
  
a.[0];;  
val it : string = "a"
```



```
let adj =  
  Array.ofList [ [1;3];  
                 [4];  
                 [4;5];  
                 [1];  
                 [3];  
                 [5]] ;;
```

```
let g6 = (6,adj);;
```

```
g6;;  
val it : int * int list []  
= (6, [[1; 3]; [4]; [4; 5]; [1]; [3]; [5]])
```






```
let (d,f,pi) = dfs(g6);;
```

d for Discovery times.  
f for finishing times.  
pi for predecessors.

```
>
```

```
val pi : int [] = [| -1; 0; -1; 4; 1; 2 |]  
val f : int [] = [| 8; 7; 12; 5; 6; 11 |]  
val d : int [] = [| 1; 2; 9; 4; 3; 10 |]
```





## A while loop

`while b do e`

can be written as the function

```
let rec wh() = if b then (e ; wh()) else ()
```

```
let mutable x = 0
```

```
while x < 10 do x<-x+1
```

```
let rec wh() = if x<10 then (x<-x+1;wh()) else ()
```

The compiler will generate the same code for the while loop and the tail-recursive `wh` function.





- F# is an excellent imperative language
- the combination of imperative and applicative constructs is powerful

In practice, only use imperative features when there is a good reason for it.



- Memory management: the stack and the heap
- Iterative (tail-recursive) functions is a simple technique to deal with efficiency in certain situations, e.g.
  - to avoid evaluations with a huge amount of **pending operations**, e.g.

$$7+(6+(5\cdots+f\ 2\cdots))$$

- to avoid inadequate use of @ in recursive declarations.
- Iterative functions with accumulating parameters correspond to while-loops
- The notion: continuations, provides a general applicable approach



Consider the following declaration:

```
let rec fact = function
  | 0 -> 1
  | n -> n * fact(n-1);;
val fact : int -> int
```

- What **resources** are needed to compute `fact(N)`?

Considerations:

- **Computation time**: number of individual computation steps.
- **Space**: the maximal memory needed during the computation to represent expressions and bindings.

## An example: Factorial function (II)



```
let rec fact = function
  | 0 -> 1
  | n -> n * fact(n-1);;
val fact : int -> int
```

Evaluation:

```
fact(N)
~> (n * fact(n-1) , [n ↦ N])
~> N * fact(N-1)
~> N * (n * fact(n-1) , [n ↦ N-1])
~> N * ((N-1) * fact(N-2))
⋮
~> N * ((N-1) * ((N-2) * (⋯ (4 * (3 * (2 * 1))) ⋯ )))
~> N * ((N-1) * ((N-2) * (⋯ (4 * (3 * 2)) ⋯ )))
⋮
~> N!
```

Time and space demands: **proportional to  $N$**       **Is this satisfactory?**



The infix operator `@` (called ‘append’) joins two lists:

$$\begin{aligned}[X_1; X_2; \dots; X_m] @ [Y_1; Y_2; \dots; Y_n] \\ = [X_1; X_2; \dots; X_m; Y_1; Y_2; \dots; Y_n]\end{aligned}$$

Properties

$$\begin{aligned}[] @ ys &= ys \\ [X_1; X_2; \dots; X_m] @ ys &= X_1 :: ([X_2; \dots; X_m] @ ys)\end{aligned}$$

Declaration

```
let rec (@) xs ys =  
  match xs with  
  | []       -> ys  
  | x::xs'   -> x::(xs' @ ys);;  
val (@) : 'a list -> 'a list -> 'a list
```

- Execution time is linear in the size of the first list

## Another example: Naive reversal (I)



```
let rec naiveRev = function
  | []      -> []
  | x::xs -> naiveRev xs @ [x];;
val naiveRev : 'a list -> 'a list
```

Evaluation of  $\text{naiveRev } [x_1, x_2, \dots, x_n]$ :

```
naiveRev [x1, x2, ..., xn]
↪ naiveRev [x2, ..., xn] @ [x1]
↪ (naiveRev [x3, ..., xn] @ [x2]) @ [x1]
⋮
↪ ((...(([] @ [xn]) @ [xn-1]) @ ... @ [x2]) @ [x1])
```

Space demands: proportional to  $n$

satisfactory

Time demands: proportional to  $n^2$

not satisfactory



**Efficient** solutions are obtained by using *more general functions*:

$$\begin{aligned}\text{factA}(n, m) &= n! \cdot m, \text{ for } n \geq 0 \\ \text{revA}([x_1, \dots, x_n], ys) &= [x_n, \dots, x_1] @ ys\end{aligned}$$

We have:

$$\begin{aligned}n! &= \text{factA}(n, 1) \\ \text{rev } [x_1, \dots, x_n] &= \text{revA}([x_1, \dots, x_n], [ ])\end{aligned}$$

*m* and *ys* are called *accumulating parameters*. They are used to hold the temporary result during the evaluation.



```
let rec factA = function
  | (0,m) -> m
  | (n,m) -> factA(n-1,n*m) ;;
```

An evaluation:

```
factA(5,1)
~> (factA(n-1,n*m), [n ↦ 5, m ↦ 1])
~> factA(4,5)
~> (factA(n-1,n*m), [n ↦ 4, m ↦ 5])
~> factA(3,20)
~> ...
~> factA(0,120) ~> (m, [m ↦ 120]) ~> 120
```

Space demand: **constant**.

Time demands: **proportional to  $n$**





```
let rec revA = function
  | ([], ys)      -> ys
  | (x::xs, ys) -> revA(xs, x::ys) ;;
```

An evaluation:

```
      revA([1,2,3], [])
  ~> revA([2,3], 1::[])
  ~> revA([3], 2::[1])
  ~> revA([3], [2,1])
  ~> revA([], 3::[2,1])
  ~> revA([], [3,2,1])
  ~> [3,2,1]
```

Space and time demands:

proportional to  $n$  (the length of the first list)



The declarations of `factA` and `revA` are *tail-recursive functions*

- the recursive call is the *last function application* to be evaluated in the body of the declaration e.g. `facA(3, 20)` and `revA([3], [2, 1])`
- only *one set* of bindings for argument identifiers is needed during the evaluation



```
let rec factA = function
  | (0,m) -> m
  | (n,m) -> factA(n-1,n*m)
              (* recursive "tail-call" *)
```

- only one set of bindings for argument identifiers is needed during the evaluation

```
factA(5,1)
~> (factA(n,m), [n ↦ 5, m ↦ 1])
~> (factA(n-1,n*m), [n ↦ 5, m ↦ 1])
~> factA(4,5)
~> (factA(n,m), [n ↦ 4, m ↦ 5])
~> (factA(n-1,n*m), [n ↦ 4, m ↦ 5])
~> ...
~> factA(0,120) ~> (m, [m ↦ 120]) ~> 120
```



```
let xs16 = List.init 1000000 (fun i -> 16);;  
val xs16 : int list = [16; 16; 16; 16; 16; ...]  
  
#time;; // a toggle in the interactive environment  
  
for i in xs16 do let _ = fact i in ();;  
Real: 00:00:00.051, CPU: 00:00:00.046, ...  
  
for i in xs16 do let _ = factA(i,1) in ();;  
Real: 00:00:00.024, CPU: 00:00:00.031, ...
```

The performance gain of `factA` is much better than the indicated factor 2 because the `for` construct alone uses about 12 ms:

```
for i in xs16 do let _ = () in ();;  
Real: 00:00:00.012, CPU: 00:00:00.015, ...
```

Real: time elapsed by the execution.    CPU: time spent by all cores.



```
let xs20000 = [1 .. 20000];;

naiveRev xs20000;;
Real: 00:00:07.624, CPU: 00:00:07.597,
GC gen0: 825, gen1: 253, gen2: 0
val it : int list = [20000; 19999; 19998; ...]

revA(xs20000, []);;
Real: 00:00:00.001, CPU: 00:00:00.000,
GC gen0: 0, gen1: 0, gen2: 0
val it : int list = [20000; 19999; 19998; ...]
```


- The naive version takes **7.624 seconds** - the iterative just **1 ms**.
- The use of append (@) has been reduced to a use of cons (: :). This has a dramatic effect of the garbage collection:
  - No object is reclaimed when `revA` is used
  - **825+253** obsolete objects were reclaimed using the naive version

Let's look at memory management



- Primitive values are allocated on the stack
- Composite values are allocated on the heap

```
let xs = [5;6;7];;  
let ys = 3::4::xs;;  
let zs = xs @ ys;;  
let n = 27;;
```





No **unnecessary copying** is done:

- 1 The linked lists for  $ys$  is not copied when building a linked list for  $y :: ys$ .
- 2 Fresh cons cells are made for the elements of  $xs$  only when building a linked list for  $xs @ ys$ .

since a list is a functional (immutable) data structure


The running time of  $@$  is linear in the length of its first argument.



Example:

```
let zs = let xs = [1;2]
         let ys = [3;4]
         xs@ys;;
```

Initial stack and heap prior to the evaluation of the local declarations:






# Operations on stack: Push



```
let zs = let xs = [1;2]
        let ys = [3;4]
        xs@ys;;
```


Evaluation of the local declarations initiated by **pushing** a new stack frame onto the stack:



The auxiliary entry **result** refers to the value of the `let`-expression.



The top stack frame is **popped** from the stack when the evaluation of the `let`-expression is completed:



The resulting heap contains two **obsolete** cells marked with '†'. They will be reclaimed by the next GC.



The memory management system uses a *garbage collector* to reclaim obsolete cells in the heap behind the scene.

The garbage collector manages the heap as partitioned into three groups or *generations*: `gen0`, `gen1` and `gen2`, according to their age. The objects in `gen0` are the youngest while the objects in `gen2` are the oldest.

The typical situation is that objects die young and the garbage collector is designed for that situation.

Example:

```
naiveRev xs20000;;  
Real: 00:00:07.624, CPU: 00:00:07.597,  
GC gen0: 825, gen1: 253, gen2: 0  
val it : int list = [20000; 19999; 19998; ...]
```



The stack is big:

```
let rec bigList n = if n=0 then [] else 1::bigList(n-1);;
bigList 120000;;
val it : int list = [1; 1; 1; 1; 1; 1; 1; 1;...]
bigList 130000;;
Process is terminated due to StackOverflowException.
```

More than  $1.2 \cdot 10^5$  stack frames are pushed in recursive calls.

Stack size depends on setup - you may have to use much larger lists.

The heap is much bigger:

```
let rec bigListA n xs = if n=0 then xs
                        else bigListA (n-1) (1::xs);;
let xsVeryBig = bigListA 12000000 [];;
val xsVeryBig : int list = [1; 1; 1; 1; 1; 1;...]
let xsTooBig = bigListA 13000000 [];;
System.OutOfMemoryException: ...
```

A list with more than  $1.2 \cdot 10^7$  elements can be created.

The iterative `bigListA` function does not exhaust the stack. **WHY?**



Tail-recursive functions are also called *iterative functions*.

- The function  $f(n, m) = (n - 1, n * m)$  is iterated during evaluations for `factA`.
- The function  $g(x :: xs, ys) = (xs, x :: ys)$  is iterated during evaluations for `revA`.

The correspondence between tail-recursive functions and while loops is established in the textbook.

An example:

```
let factW n =  
  let ni = ref n  
  let r  = ref 1  
  while !ni > 0 do  
    r := !r * !ni ; ni := !ni - 1  
  !r;;
```



A function  $g : \tau \rightarrow \tau'$  is an *iteration of*  $f : \tau \rightarrow \tau$  if it is an instance of:

```
let rec g z = if p z then g(f z) else h z
```

for suitable predicate  $p : \tau \rightarrow \text{bool}$  and function  $h : \tau \rightarrow \tau'$ .

The function  $g$  is called an *iterative (or tail-recursive) function*.

Examples: `factA` and `revA` are easily declared in the above form:

```
let rec factA(n,m) =  
  if n <> 0 then factA(n-1,n*m) else m;;
```

```
let rec revA(xs,ys) =  
  if not (List.isEmpty xs)  
  then revA(List.tail xs, (List.head xs)::ys)  
  else ys;;
```



Consider: `let rec g z = if p z then g (f z) else h z`

Evaluation of the `g v`:

```
g v
~> (if p z then g (f z) else h z, [z ↦ v])
~> (g (f z), [z ↦ v])
~> g (f1 v)
~> (if p z then g (f z) else h z, [z ↦ f1 v])
~> (g (f z), [z ↦ f1 v])
~> g (f2 v)
~> ...
~> (if p z then g (f z) else h z, [z ↦ fn v])
~> (h z, [z ↦ fn v])
~> h (fn v)
```

suppose  $p(f^n v) \rightsquigarrow \text{false}$



Observe two desirable properties:

- there are  $n$  recursive calls of  $g$ ,
- at most *one binding* for the argument pattern  $z$  is 'active' at any stage in the evaluation, and
- the iterative functions require *one* stack frame only.





Iterative functions are executed efficiently:

```
#time;;
```

```
for i in 1 .. 1000000 do let _ = factA(16,1) in ();;  
Real: 00:00:00.024, CPU: 00:00:00.031,  
GC gen0: 0, gen1: 0, gen2: 0  
val it : unit = ()
```

```
for i in 1 .. 1000000 do let _ = factW 16 in ();;  
Real: 00:00:00.048, CPU: 00:00:00.046,  
GC gen0: 9, gen1: 0, gen2: 0  
val it : unit = ()
```

- the tail-recursive function actually is faster than the imperative while-loop based version

## Example: Fibonacci numbers (I)



A declaration based directly on the mathematical definition:

```
let rec fib = function
  | 0 -> 0
  | 1 -> 1
  | n -> fib(n-1) + fib(n-2);;
val fib : int -> int
```

is highly inefficient. For example:

```
fib 4
~> fib 3 + fib 2
~> (fib 2 + fib 1) + fib 2
~> ((fib 1 + fib 0) + fib 1) + fib 2
~> ... ~> 2 + (fib 1 + fib 0)
~> ...
```

Ex: `fib 44` requires around  $10^9$  evaluations of base cases.

## Example: Fibonacci numbers (II)



An iterative solution gives high efficiency:

```
fun rec itfib(n,a,b) = if n <> 0
                        then itfib(n-1,a+b,a)
                        else a;;
```

The expression `itfib( $n$ , 0, 1)` evaluates to  $F_n$ , for any  $n \geq 0$ :

- Case  $n = 0$ : `itfib(0, 0, 1)  $\rightsquigarrow$  0 ( $= F_0$ )`
- Case  $n > 0$ :

```
    itfib(n, 0, 1)
 $\rightsquigarrow$  itfib(n-1, 1, 0) = itfib(n-1,  $F_1$ ,  $F_0$ )
 $\rightsquigarrow$  itfib(n-2,  $F_1 + F_0$ ,  $F_1$ )
 $\rightsquigarrow$  itfib(n-2,  $F_2$ ,  $F_1$ )
    ⋮
 $\rightsquigarrow$  itfib(0,  $F_n$ ,  $F_{n-1}$ )
 $\rightsquigarrow$   $F_n$ 
```



Accumulating parameters are not sufficient to achieve a tail-recursive version for arbitrary recursive functions.

Consider for example:

```
type BinTree<'a> =  
  | Leaf  
  | Node of BinTree<'a> * 'a * BinTree<'a>;;  
  
let rec count = function  
  | Leaf          -> 0  
  | Node(tl,n,tr) -> count tl + count tr + 1;;
```

A counting function:

```
countA: int -> BinTree<'a> -> int
```

using an accumulating parameter will **not be tail-recursive** due to the expression containing recursive calls on the left and right sub-trees.  
(Ex. 9.8)



**Continuation:** A function for the “rest” of the computation.

The continuation-based version of `bigList` has a continuation

```
c: int list -> int list
```

as argument:

```
let rec bigListC n c =  
  if n=0 then c []  
  else bigListC (n-1) (fun res -> c(1::res));;  
val bigListC : int -> (int list -> 'a) -> 'a
```

- Base case: “feed” the result of `bigList 0` into the continuation `c`.
- Recursive case: let `res` denote the value of `bigList (n-1)`:
  - The rest of the computation of `bigList n` is `1::res`.
  - The continuation of `bigListC (n-1)` is  

```
fun res -> c(1::res)
```



- `bigListC` is a tail-recursive function, and
- the calls of `c` are tail calls in the base case of `bigListC` and in the continuation: `fun res -> c(1::res)`.

The stack will hence neither grow due to the evaluation of recursive calls of `bigListC` nor due to calls of the continuations that have been built in the heap:

```
bigListC 16000000 id;;  
Real: 00:00:08.586, CPU: 00:00:08.314,  
GC gen0: 80, gen1: 60, gen2: 3  
val it : int list = [1; 1; 1; 1; 1; ...]
```

- Slower than `bigList`
- Can generate longer lists than `bigList`

## Example: Tail-recursive count



```
let rec countC t c =  
  match t with  
  | Leaf          -> c 0  
  | Node(tl,n,tr) ->  
    countC tl (fun vl -> countC tr (fun vr -> c(vl+vr+1)))  
val countC : BinTree<'a> -> (int -> 'b) -> 'b  
  
countC (Node(Node(Leaf,1,Leaf),2,Node(Leaf,3,Leaf))) id;  
val it : int = 3
```

- Both calls of `countC` are tail calls
- The calls of the `c` is tail call

Hence, the stack will not grow when evaluating `countC t c`.

- `countC` can handle bigger trees than `count`
- `count` is faster



- Loops in imperative languages corresponds to a *special case* of recursive function called tail recursive functions.
- Have iterative functions in mind when dealing with efficiency, e.g.
  - to avoid evaluations with a huge amount of pending operations
  - to avoid inadequate use of @ in recursive declarations.
- Memory management: stack, heap, garbage collection
- Continuations – provide a technique to turn arbitrary recursive functions into tail-recursive ones.

trades stack for heap

Note: Iterative function does not replace algorithmic idea and the use of good algorithms and datastructure.